

**Lab 10 · Floating Point Calculation***Lecturer: Philippos Mordohai, Shudong Hao**Date: November 9***Contents**

<b>1 Objective</b>	<b>1</b>
<b>2 Task</b>	<b>1</b>
<b>3 Requirements</b>	<b>2</b>
<b>Appendix A Architecture</b>	<b>3</b>
<b>Appendix B Basic Instructions</b>	<b>3</b>
B.1 Arithmetic . . . . .	3
B.2 Moving Real Numbers . . . . .	4
B.3 Converting Precisions . . . . .	4
<b>Appendix C Printing Using printf</b>	<b>4</b>
<b>Appendix D Debugging</b>	<b>6</b>

**1 Objective**

Get to know and use floating point numbers in ARMv8 assembly.

**2 Task**

Given a math function  $y = f(x)$  and an interval  $[a, b]$ , use the rectangle rule to approximate the integral.

Your data segment should contain at least these variables:

```

1 .data
2 a:    .double ... // left limit
3 b:    .double ... // right limit

```

```

4  n:      .double ... // number of rectangles under the curve
5
6  .bss
7  result: .skip 8     // the result of integral

```

Using the example shown in Figure 1, the data segment should look like this:

```

1  .data
2  a:      .double 0.0 // left limit = 0
3  b:      .double 2.0 // right limit = 2
4  n:      .double 10 // 10 rectangles
5  .bss
6  result: .skip 8     // the result of integral

```




Figure 1: Example of approximation.

You can assume the curve, at least the curve between  $[a, b]$  is always above  $x$ -axis.

### 3 Requirements

- The math function you're going to use is  $y = 2.5x^3 - 15.5x^2 + 20x + 15$ , and the integral is between  $[-0.5, 5]$ ;
- The specific number of rectangles is up to you, as long as your calculation is correct, and the difference between the approximation and the actual integral is no larger than  $1e - 3$ ;

- Please use `printf` to print out: your approximation result, the actual value of the integral (ok to calculate by hand), and the difference between them;
- You'll get zero point if you only calculated based on the integral formula instead of approximation.

### What to Submit

A single assembly source code file `.s`.

## A Architecture

Real numbers are a different species than integer numbers: they're encoded in a different way (if you're curious, check out Chapter 3); they are stored in a separate group of registers, instead of `X0 – X31`; they are calculated in a unit called FPU (floating point unit), not ALU (!); and they use a totally different set of instructions.

**Registers.** Real numbers in ARM have 5 types of precision, but we mostly use 2 of them: single (encoded in 32 bits), and double (64 bits). Single precision numbers correspond to `float` type in C, and double precision numbers are, well, `double` type. Therefore, correspondingly, the registers that hold them are `Hn` (half precision), `Sn` (single), and `Dn` (double), where `n` ranging from 0 to 31 is the register number.

**Usage.** Like integer registers, we use `D0 – D7` to pass parameters to procedures as well as store return values. `D8 – D15` are preserved across calls, and `D16 – D31` can be used as temporary registers. The same applies to single and half precision registers.

## B Basic Instructions

For real numbers, they have a separate set of instructions, but thankfully they are very similar to the ones we know so far.

### B.1 Arithmetic

Typically, when using floating point registers, the instruction has an `F` prefix. For example:

```
1 FADD S0, S2, S3 // S0 = S2 + S3
2 FADD D19, D12, D12 // D19 = D12 + D12
```

However, load and store are the same as before.

## B.2 Moving Real Numbers

You can move an immediate real number to an S or D register using `FMOV`, but there's a restriction on that. Based on ARMv8 manual, only numbers that can be expressed as  $\pm \frac{n}{16} \times 2^r$  where  $n \in [16, 31]$  and  $r \in [-3, 4]$  can be moved. Numbers such as `x.0` (integers), `x.5`, `x.25` are fine.

Therefore, we recommend that you just store all real numbers in the `data` segment, and load them into registers, and use `FMOV` between registers. Assume we have a number declared as such:

```
1 .data
2 pi: .float 3.1415
```

Then to move this number into a register, we should do:

```
1 ADR    X0, pi
2 LDUR   S0, [X0]
3 FMOV   S1, S0
```

## B.3 Converting Precisions

Sometimes we'd like to cast a floating point to a double precision number. We can't just `FMOV` an S register to a D register. The instruction we need to use is `FCVT`:

```
1 FCVT    D0, S0    // Upcast, gain precision
2 FCVT    S1, D1    // Downcast, lose precision
3 SCVTF   D2, X2    // Convert an integer to a real number
4 FCVTZS  X3, D3    // Convert a real number to an integer (the fraction part_
    ↳is removed)
```

For other instructions, the best resource out there is ARM64's reference sheet: <https://developer.arm.com/documentation/100076/0100/a64-instruction-set-reference/a64-floating-point-instructions>.

## C Printing Using printf

Printing is a little bit tricky. As usual, we need to load a string to `X0`, but the rest of the parameters are passed from `D0` – `D7`.

```

1 .data
2 fmt_str: .ascii "%lf %lf\n\0"
3 number1: .double 3.1415
4 number2: .double 10
5 ...
6 ADR    X0, fmt_str    // Load address of the string
7 ADR    X1, number1    // Load address of number1
8 LDUR   D0, [X1]       // Load number1 to D0
9 ADR    X1, number2
10 LDUR   D1, [X1]
11 BL     printf

```

When `printf` recognizes `%lf`, it doesn't go to `X1` to fetch the number; instead it goes to `D0`.

If you use `printf` to print both integer values (`int`, `char`, `long int`) and floating point values, move corresponding numbers into their registers in order:

```

1 .data
2 fmt_str: .ascii "%d = %lf, %d = %lf\n\0"
3 ...
4 ADR    X0, fmt_str
5 LDUR   X1, [...]     // integer #1
6 LDUR   D0, [...]     // floating point #1
7 LDUR   X2, [...]     // integer #2
8 LDUR   D1, [...]     // floating point #2

```

**!Caution!** In case you declare your number as `.float` like following:

```

1 .data
2 fmt_str: .ascii "%f\n\0"
3 fpnum:   .float 3.14

```

you would need to cast `fpnum` from `float` to `double` because `printf` doesn't check `S` registers at all, and automatically uses double precision as output format (even if your format is `%f` instead of `%lf`):

```

1 ADR    X0, fmt_str
2 ADR    X1, fpnum
3 LDUR   S0, [X1]

```

```
4 FCVT    D0, S0
5 BL      printf
```

And if you declare your number as `.float`, you cannot load it directly into D registers.

Thus, the easiest way to avoid those situations is just declare your number as `.double`, and load it into D registers all the time.

## D Debugging

The floating point registers cannot be viewed in the register panel, so we would have to use `p` to print out:

```
1 p/f $d0
```

where `/f` is to print out the register value as a floating point.